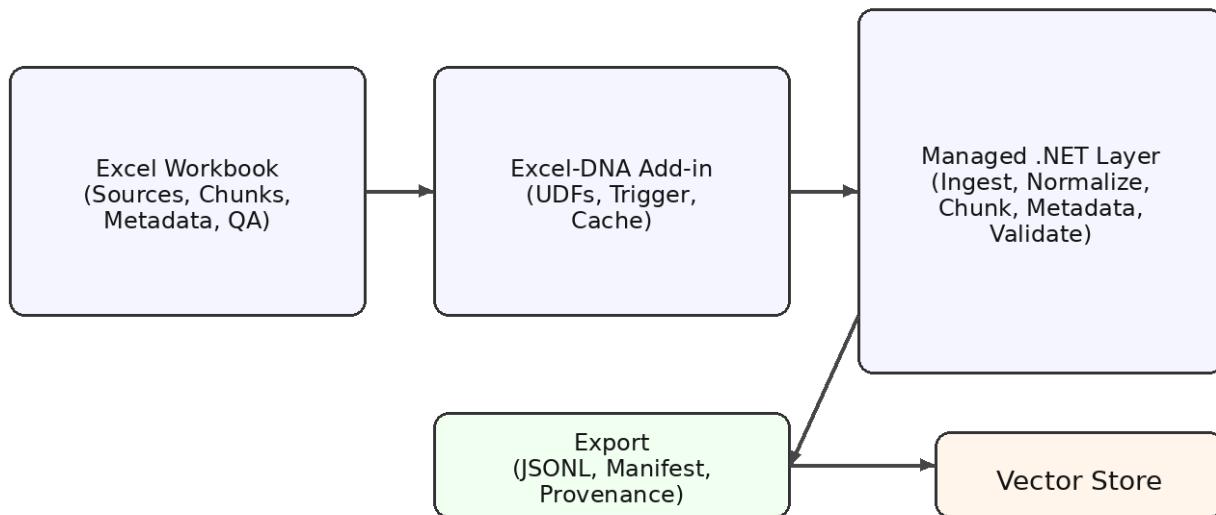


# ExcelRAG

Turn Excel into a first-class workspace for preparing Retrieval-Augmented Generation (RAG) data — natively, transparently, and at scale.

## Vision

ExcelRAG Pipeline Overview



*Excel as a visual, auditable RAG preparation pipeline*

*Diagram note: data flows from Excel tables → Excel-DNA UDFs → managed processing → export packages, keeping every intermediate artifact visible in the worksheet.*

**ExcelRAG** is an Excel-DNA add-in that transforms Excel into a powerful, trustworthy workspace for preparing data used in RAG pipelines. It brings modern data prep and knowledge-curation workflows into Excel while keeping the experience formula-driven, transparent, and fully inspectable.

The guiding idea:

*If someone already trusts Excel to prepare business data, they should be able to prepare RAG-ready knowledge with the same mental model.*

ExcelRAG is **not** a production RAG engine. It is a **data preparation, governance, and export layer** built for:

- structured content ingestion
  - chunking and normalization
  - metadata enrichment
  - quality checks and audits
  - vector-ready export packages
- 

## Why Excel for RAG?

Excel already excels at:

- tabular structure and auditability
- formulas for transformation and enrichment
- versioned snapshots of curated knowledge
- collaborative review and approval
- quick visualization and QA checks

What it lacks is RAG-specific tooling:

- document ingestion and parsing
- chunking strategies
- metadata extraction
- deduplication and similarity checks
- vector packaging and export

ExcelRAG fills that gap using **Excel-native UDFs** and a small, explicit processing engine behind the scenes.

---

## Core Design Principles

### 1) Excel remains Excel

- No VBA macros
- No external Python runtime
- No opaque background automation

Every capability is exposed as **deterministic Excel functions** that are visible, reviewable, and easy to test.

---

## 2) Explicit, user-controlled processing

ExcelRAG avoids surprise side effects by requiring a user-controlled **trigger** for any operation that mutates state or performs heavy processing.

```
=RAG.INGEST(source, options, trigger)  
=RAG.CHUNK(content, strategy, options, trigger)
```

This matches Excel's recalculation model and keeps data prep deliberate and auditable.

---

## 3) Inspectable, cell-level transparency

Every step produces tangible, inspectable outputs in the grid:

- parsed text and tables
- chunk boundaries
- metadata fields
- token counts
- validation warnings

If it can be used for RAG, it can be **seen and verified** in Excel.

---

## 4) RAG-first data quality

ExcelRAG is designed for **clean, consistent, and richly annotated** content. It prioritizes:

- normalization
  - traceability to source
  - metadata completeness
  - data provenance
  - export-ready structure
- 

## RAG Workflow Support

ExcelRAG supports a typical RAG preparation lifecycle:

### 1. Ingest

- PDFs, DOCX, HTML, Markdown, CSV, URLs
- Text extraction with source metadata

## **2. Normalize & Clean**

- whitespace and formatting normalization
- boilerplate removal
- table flattening

## **3. Chunk**

- sentence-aware or token-based chunking
- overlap control
- section-aware chunking for headings

## **4. Enrich Metadata**

- document titles and sections
- author/owner fields
- business tags
- PII flags or sensitivity levels

## **5. Quality Checks**

- duplicate detection
- minimum/maximum token constraints
- missing metadata checks
- hallucination risk flags (e.g., low coverage)

## **6. Package & Export**

- vector-store friendly JSONL
- embeddings-ready manifests
- provenance maps

---

# **Key Features (Planned Add-in Functions)**

## **Ingestion & Parsing**

- `RAG . INGEST` — import files or URLs with parsing options
- `RAG . EXTRACT_TEXT` — extract text from common document formats
- `RAG . EXTRACT_TABLES` — flatten tables to text or structured ranges

## **Chunking & Structuring**

- `RAG . CHUNK` — generate chunk lists using strategies (sentence, token, heading)
- `RAG . CHUNK_STATS` — token counts, overlap metrics, and coverage
- `RAG . SECTIONS` — parse headings and structure long docs

## Metadata & Enrichment

- RAG.METADATA — create consistent metadata records
- RAG.TAG — generate or validate tags against a taxonomy
- RAG.PROVENANCE — map chunks back to source offsets

## Quality & Governance

- RAG.DEDUP — detect duplicate or near-duplicate chunks
- RAG.VALIDATE — rule-based data quality checks
- RAG.RISK\_FLAGS — highlight content risk or low-coverage areas

## Export & Packaging

- RAG.EXPORT\_JSONL — produce vector-ready export files
- RAG.EXPORT\_MANIFEST — embeddings manifests with metadata and source links
- RAG.EXPORT\_TRACE — end-to-end traceability for governance

---

## Ribbon (Planned UI)

ExcelRAG includes a dedicated **RAG** ribbon tab for guided workflows and governance-friendly actions. The ribbon complements UDFs by providing quick actions, templates, and validation panels without hiding any outputs from the worksheet.

### Tab: RAG

- **Ingest**
  - Import Files (PDF, DOCX, HTML, MD)
  - Import URL
  - Manage Sources
- **Normalize**
  - Clean Text
  - Remove Boilerplate
  - Flatten Tables
- **Chunk**
  - Sentence-Aware Chunking
  - Token-Based Chunking
  - Preview Chunk Boundaries
- **Metadata**

- Apply Metadata Template
- Tag Validation
- Provenance Mapping
- **Quality**
  - Duplicate Scan
  - Coverage Check
  - Missing Metadata Report
- **Export**
  - Export JSONL
  - Export Manifest
  - Open Export Folder

- **Diagnostics**

- Processing Log
- Cache Inspector
- Settings

The ribbon actions generate or refresh **worksheet tables** so every step remains visible and reviewable.

---

## Example Dataset and UDF Effects

Below is a small example dataset representing three source documents in a worksheet table.

### **Input: Sources**

SourceId	Title	Content	Owner	Tags
DOC-001	Employee Onboarding	Welcome to Contoso. New hires must complete security training within 10 business days.	HR	onboarding;policy
DOC-002	VPN Access	VPN access is required for remote work. Submit the VPN request form and complete MFA setup.	IT	security;remote
DOC-003	Expense Policy	Receipts are required for expenses over <b>25. Mileage is reimbursed at 0.67 per mile.</b>	Finance	policy;travel

**Effect: RAG.CHUNK** (sentence-aware, max 40 tokens, overlap 8)

ChunkId	SourceId	ChunkText	TokenCount	OffsetStart	OffsetEnd
CH-0001	DOC-001	Welcome to Contoso. New hires must complete security training within 10 business days.	20	0	97
CH-0002	DOC-002	VPN access is required for remote work. Submit the VPN request form and complete MFA setup.	23	0	101
CH-0003	DOC-003	Receipts are required for expenses over <b>25. Mileage is reimbursed at 0.67 per mile.</b>	22	0	99

**Effect: RAG.METADATA** (template with SourceId, Title, Owner, Tags, Section=General)

ChunkId	SourceId	Title	Owner	Tags	Section
CH-0001	DOC-001	Employee Onboarding	HR	onboarding;policy	General
CH-0002	DOC-002	VPN Access	IT	security;remote	General
CH-0003	DOC-003	Expense Policy	Finance	policy;travel	General

**Effect: RAG.DEDUP** (near-duplicate scan, threshold 0.92)

ChunkId	DuplicateGroup	Similarity	Status
CH-0001	G-01	0.18	Unique
CH-0002	G-02	0.22	Unique
CH-0003	G-03	0.20	Unique

**Effect: RAG.VALIDATE** (rules: minTokens=10, maxTokens=250, tagsRequired=true)

ChunkId	Rule	Result	Notes
CH-0001	TagsRequired	Pass	Tags present
CH-0001	TokenRange	Pass	20 tokens
CH-0002	TagsRequired	Pass	Tags present
CH-0002	TokenRange	Pass	23 tokens
CH-0003	TagsRequired	Pass	Tags present

ChunkId	Rule	Result	Notes
CH-0003	TokenRange	Pass	22 tokens

These tables are meant to illustrate what appears directly in Excel as outputs of key UDFs, enabling quick review before export.

---

## Architecture Overview

### Excel

- └ Worksheets (tables, named ranges, checks)
- └ UDFs (RAG.INGEST, RAG.CHUNK, ...)

### Excel-DNA Add-in (.xll)

- └ Async-safe UDF execution
- └ Trigger-aware processing
- └ Local cache & registry

### Managed .NET Layer

- └ Parsing, chunking, metadata
- └ Content normalization
- └ Validation rules engine

### Export Layer

- └ JSONL, manifest, provenance
- └ Vector-store ready packages

All logic runs **in-process** with Excel for visibility and speed, while keeping data in the workbook for auditability.

---

## Intended Audience

ExcelRAG is for:

- analysts curating enterprise knowledge
- teams preparing documents for search and RAG
- compliance-minded organizations
- Excel power users who want a transparent, reviewable pipeline

It is **not** for:

- full-scale production inference
  - large-scale embedding compute
  - real-time RAG orchestration
- 

## Status

This project is **experimental and exploratory**.

The goal is to answer one question:

*What would RAG data preparation look like if it were designed for transparency and governance first?*

---

## See Also

- [ExcelDNA\\_Best\\_Practices\\_and\\_Architectural\\_Patterns\\_2026.md](#)
- [ExcelD1Playground/README.md](#)